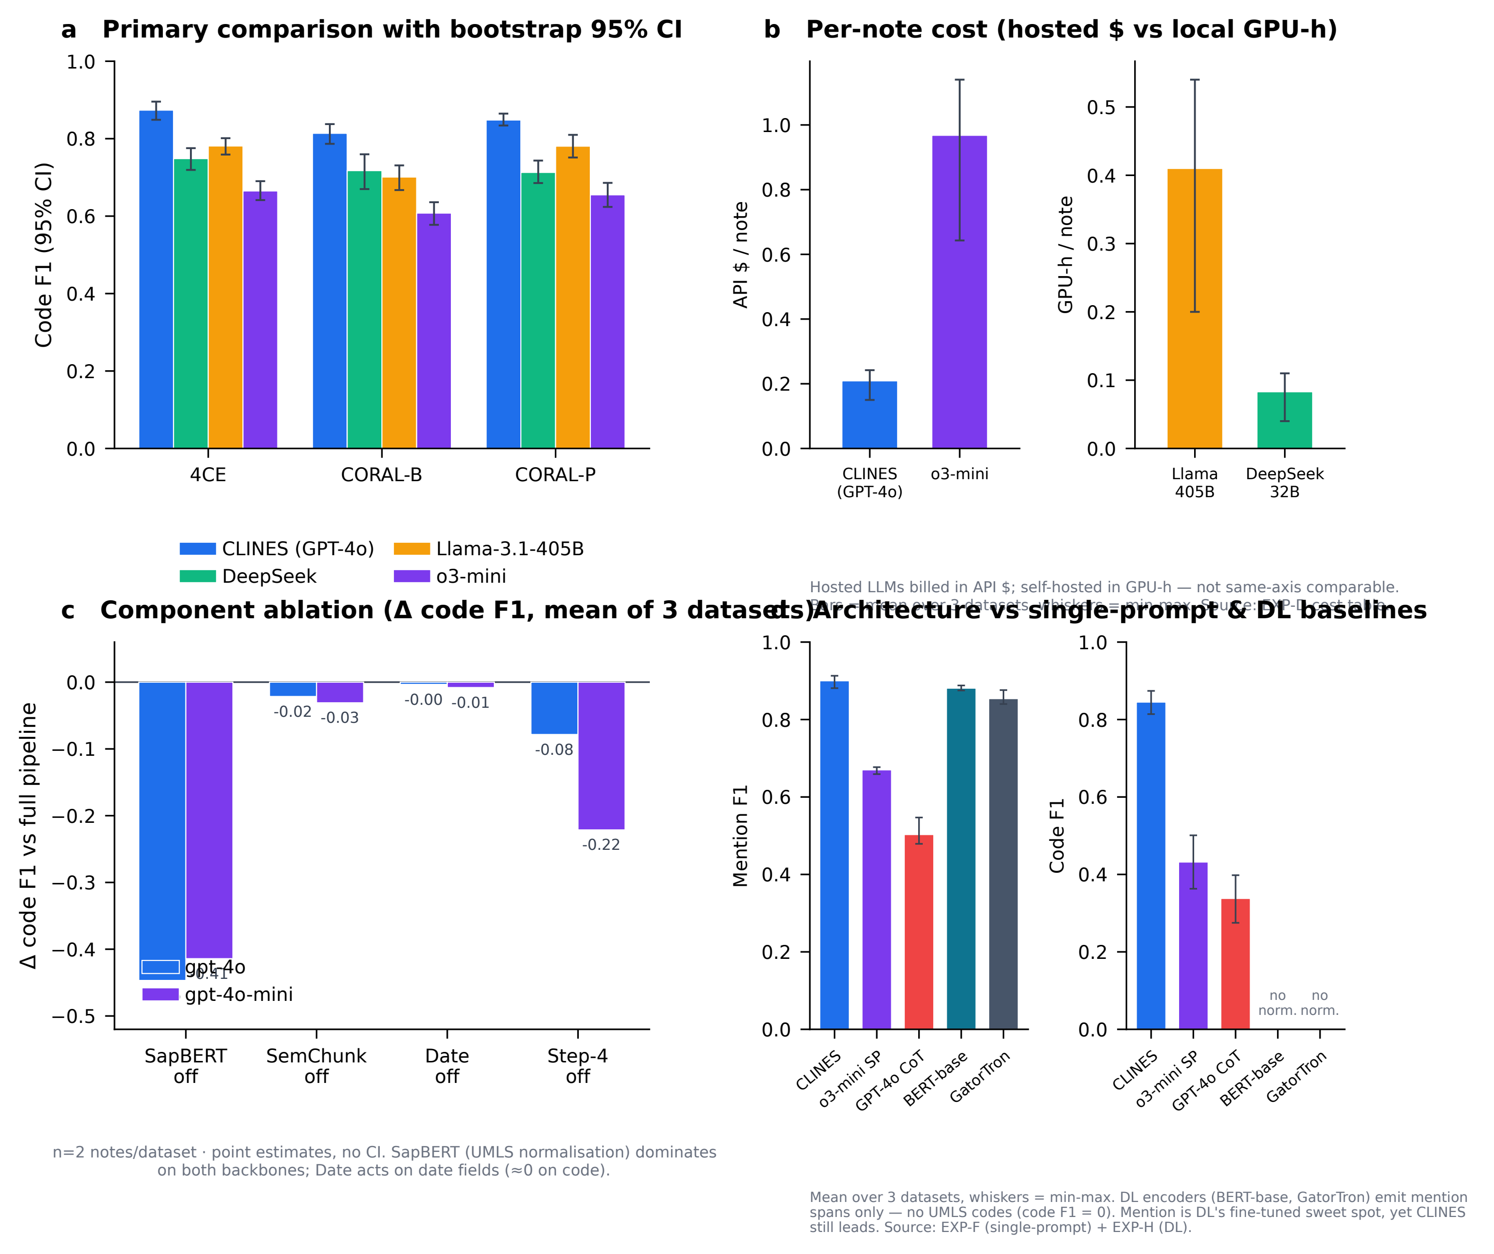

a Primary comparison with bootstrap 95% CI
Code F1 (95% CI)
1.0
0.8
0.6
0.4
0.2
0.0
4CE
CORAL-B
CORAL-P
CLINES (GPT-4o)
Llama-3.1-405B
DeepSeek
o3-mini
b Per-note cost (hosted $ vs local GPU-h)
API $ / note
1.0
0.8
0.6
0.4
0.2
0.0
CLINES (GPT-4o)
o3-mini
GPU-h / note
0.5
0.4
0.3
0.2
0.1
0.0
Llama 405B
DeepSeek 32B
Hosted LLMs billed in API $; self-hosted in GPU-h — not same-axis comparable.
Bars = mean over 3 datasets, whiskers = min-max. Source: EXP-D cost table.
c Component ablation (Δ code F1, mean of 3 datasets)
Δ code F1 vs full pipeline
0.0
−0.1
−0.2
−0.3
−0.4
−0.5
-0.02
-0.03
-0.00
-0.01
-0.08
-0.22
-0.41
gpt-4o
gpt-4o-mini
SapBERT off
SemChunk off
Date off
Step-4 off
n=2 notes/dataset · point estimates, no CI. SapBERT (UMLS normalisation) dominates on both backbones; Date acts on date fields (≈0 on code).
d Architecture vs single-prompt & DL baselines
Mention F1
1.0
0.8
0.6
0.4
0.2
0.0
CLINES
o3-mini SP
GPT-4o CoT
BERT-base
GatorTron
Code F1
1.0
0.8
0.6
0.4
0.2
0.0
no norm.
no norm.
CLINES
o3-mini SP
GPT-4o CoT
BERT-base
GatorTron
Mean over 3 datasets, whiskers = min-max. DL encoders (BERT-base, GatorTron) emit mention spans only — no UMLS codes (code F1 = 0). Mention is DL's fine-tuned sweet spot, yet CLINES still leads. Source: EXP-F (single-prompt) + EXP-H (DL).