






a Primary comparison with bootstrap 95% CI**b Per-note cost (hosted \$ vs local GPU-h)****c Component ablation (Δ code F1, mean of 3 datasets) Architecture vs single-prompt & DL baselines**

Hosted LLMs billed in API \$; self-hosted in GPU-h — not same-axis comparable.
 Mean over 3 datasets, whiskers = min-max. DL encoders (BERT-base, GatorTron) emit mention spans only — no UMLS codes (code F1 = 0). Mention is DL's fine-tuned sweet spot, yet CLINES still leads. Source: EXP-F (single-prompt) + EXP-H (DL).



n=2 notes/dataset · point estimates, no CI. SapBERT (UMLS normalisation) dominates on both backbones; Date acts on date fields (≈ 0 on code).

Mean over 3 datasets, whiskers = min-max. DL encoders (BERT-base, GatorTron) emit mention spans only — no UMLS codes (code F1 = 0). Mention is DL's fine-tuned sweet spot, yet CLINES still leads. Source: EXP-F (single-prompt) + EXP-H (DL).